MINI PROJECT

On

BOSTON HOUSING DATA

(Using R)

Prepared By:

Wyendrila Roy

http://in.linkedin.com/pub/wyendrila-roy/5/3a/876

Understanding the Data

Variable Name	Definition
MEDV	median value in $1000
DIS	distance to employment centers
CRIM	per capita crime rate
RAD	accessibility to radial highways
ZN	% land zoned for lots
INDUS	% nonretail business
TAX	property tax/$10,000
CHAS	1 on Charles River, 0 else
PT	pupil/teacher ratio
NOX	nitrogen oxide conc. (p.p.109)
Black	(% black - 63)2/10
RM	average number of rooms
LSTAT	% lower-status pop.
AGE	% built before 1940

Structure of the Data

> data(Boston, package="MASS") (This function reads the Boston dataset from the MASS Package in R)

> names(Boston) (This function gives the variable names in the Boston dataset, the definition has been provided at the beginning for reference)

```
 [1] "crim"    "zn"      "indus"   "chas"    "nox"     "rm"      "age"
 [8] "dis"     "rad"     "tax"     "ptratio" "black"   "lstat"   "medv"
```

> str(Boston) (This function gives the structure of the data, the total number of variables and observations in the data and the datatype of each variable)

```
'data.frame':       506 obs. of  14 variables:
 $ crim   : num  0.00632 0.02731 0.02729 0.03237 0.06905 ...
 $ zn     : num  18 0 0 0 0 0 12.5 12.5 12.5 12.5 ...
 $ indus  : num  2.31 7.07 7.07 2.18 2.18 2.18 7.87 7.87 7.87 7.87 ...
 $ chas   : int  0 0 0 0 0 0 0 0 0 0 ...
 $ nox    : num  0.538 0.469 0.469 0.458 0.458 0.458 0.524 0.524 0.524 0.524 ...
 $ rm     : num  6.58 6.42 7.18 7 7.15 ...
 $ age    : num  65.2 78.9 61.1 45.8 54.2 58.7 66.6 96.1 100 85.9 ...
 $ dis    : num  4.09 4.97 4.97 6.06 6.06 ...
 $ rad    : int  1 2 2 3 3 3 5 5 5 5 ...
 $ tax    : num  296 242 242 222 222 222 311 311 311 311 ...
```

$ ptratio: num 15.3 17.8 17.8 18.7 18.7 18.7 15.2 15.2 15.2 15.2 ...
$ black : num 397 397 393 395 397 ...
$ lstat : num 4.98 9.14 4.03 2.94 5.33 ...
$ medv : num 24 21.6 34.7 33.4 36.2 28.7 22.9 27.1 16.5 18.9 ...

> head(Boston) (To view the first few rows in the data in order to check the data as a sample)

	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	black	lstat	medv
1	0.00632	18	2.31	0	0.538	6.575	65.2	4.0900	1	296	15.3	396.90	4.98	24.0
2	0.02731	0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8	396.90	9.14	21.6
3	0.02729	0	7.07	0	0.469	7.185	61.1	4.9671	2	242	17.8	392.83	4.03	34.7
4	0.03237	0	2.18	0	0.458	6.998	45.8	6.0622	3	222	18.7	394.63	2.94	33.4
5	0.06905	0	2.18	0	0.458	7.147	54.2	6.0622	3	222	18.7	396.90	5.33	36.2
6	0.02985	0	2.18	0	0.458	6.430	58.7	6.0622	3	222	18.7	394.12	5.21	28.7

> summary(Boston) (To understand the basic statistics of each individual variable)

crim	zn	indus	chas
Min. : 0.00632	Min. : 0.00	Min. : 0.46	Min. :0.00000
1st Qu.: 0.08204	1st Qu.: 0.00	1st Qu.: 5.19	1st Qu.:0.00000
Median : 0.25651	Median : 0.00	Median : 9.69	Median :0.00000
Mean : 3.61352	Mean : 11.36	Mean :11.14	Mean :0.06917
3rd Qu.: 3.67708	3rd Qu.: 12.50	3rd Qu.:18.10	3rd Qu.:0.00000
Max. :88.97620	Max. :100.00	Max. :27.74	Max. :1.00000

nox	rm	age	dis
Min. :0.3850	Min. :3.561	Min. : 2.90	Min. : 1.130
1st Qu.:0.4490	1st Qu.:5.886	1st Qu.: 45.02	1st Qu.: 2.100
Median :0.5380	Median :6.208	Median : 77.50	Median : 3.207
Mean :0.5547	Mean :6.285	Mean : 68.57	Mean : 3.795
3rd Qu.:0.6240	3rd Qu.:6.623	3rd Qu.: 94.08	3rd Qu.: 5.188
Max. :0.8710	Max. :8.780	Max. :100.00	Max. :12.127

rad	tax	ptratio	black
Min. : 1.000	Min. :187.0	Min. :12.60	Min. : 0.32
1st Qu.: 4.000	1st Qu.:279.0	1st Qu.:17.40	1st Qu.:375.38
Median : 5.000	Median :330.0	Median :19.05	Median :391.44
Mean : 9.549	Mean :408.2	Mean :18.46	Mean :356.67
3rd Qu.:24.000	3rd Qu.:666.0	3rd Qu.:20.20	3rd Qu.:396.23
Max. :24.000	Max. :711.0	Max. :22.00	Max. :396.90

lstat	medv

Min. : 1.73	Min. : 5.00
1st Qu.: 6.95	1st Qu.:17.02
Median :11.36	Median :21.20
Mean :12.65	Mean :22.53
3rd Qu.:16.95	3rd Qu.:25.00
Max. :37.97	Max. :50.00

Viewing the median value of owner occupied homes in $1000 by using histogram and normal Q-Q Plot

```
attach(Boston)
hist(medv)
qqnorm(medv)
qqline(medv)
```

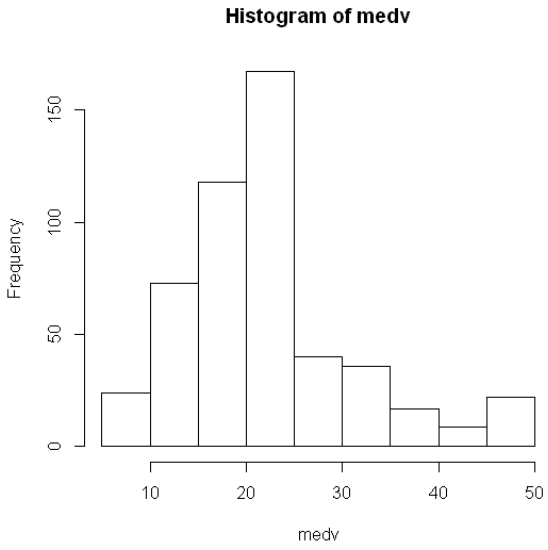

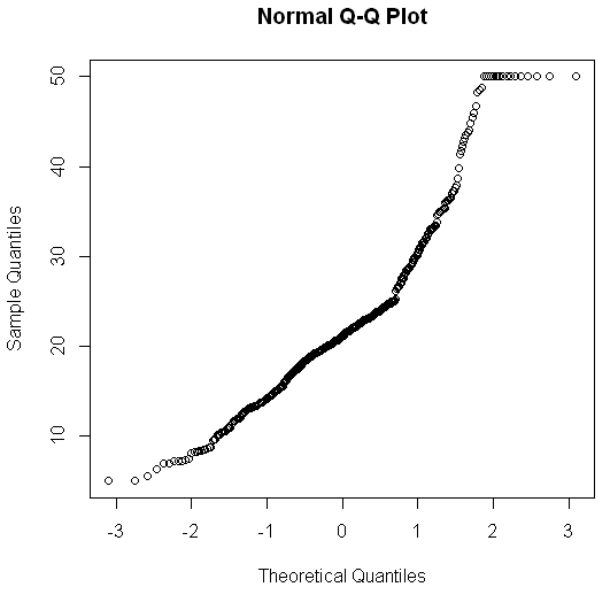

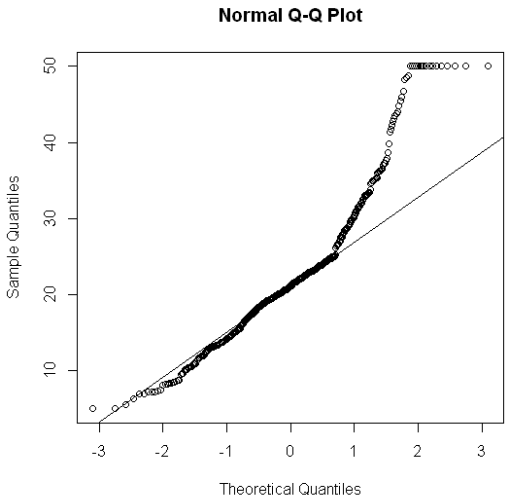


```
> cor(Boston)
```
(To understand the correlation between each individual variable)

	crim	zn	indus	chas	nox
crim	1.00000000	-0.20046922	0.40658341	-0.055891582	0.42097171
zn	-0.20046922	1.00000000	-0.53382819	-0.042696719	-0.51660371
indus	0.40658341	-0.53382819	1.00000000	0.062938027	0.76365145
chas	-0.05589158	-0.04269672	0.06293803	1.000000000	0.09120281
nox	0.42097171	-0.51660371	0.76365145	0.091202807	1.00000000
rm	-0.21924670	0.31199059	-0.39167585	0.091251225	-0.30218819
age	0.35273425	-0.56953734	0.64477851	0.086517774	0.73147010
dis	-0.37967009	0.66440822	-0.70802699	-0.099175780	-0.76923011
rad	0.62550515	-0.31194783	0.59512927	-0.007368241	0.61144056
tax	0.58276431	-0.31456332	0.72076018	-0.035586518	0.66802320
ptratio	0.28994558	-0.39167855	0.38324756	-0.121515174	0.18893268
black	-0.38506394	0.17552032	-0.35697654	0.048788485	-0.38005064
lstat	0.45562148	-0.41299457	0.60379972	-0.053929298	0.59087892
medv	-0.38830461	0.36044534	-0.48372516	0.175260177	-0.42732077

	rm	age	dis	rad	tax	ptratio
crim	-0.21924670	0.35273425	-0.37967009	0.625505145	0.58276431	0.2899456
zn	0.31199059	-0.56953734	0.66440822	-0.311947826	-0.31456332	-0.3916785
indus	-0.39167585	0.64477851	-0.70802699	0.595129275	0.72076018	0.3832476
chas	0.09125123	0.08651777	-0.09917578	-0.007368241	-0.03558652	-0.1215152
nox	-0.30218819	0.73147010	-0.76923011	0.611440563	0.66802320	0.1889327
rm	1.00000000	-0.24026493	0.20524621	-0.209846668	-0.29204783	-0.3555015
age	-0.24026493	1.00000000	-0.74788054	0.456022452	0.50645559	0.2615150
dis	0.20524621	-0.74788054	1.00000000	-0.494587930	-0.53443158	-0.2324705
rad	-0.20984667	0.45602245	-0.49458793	1.000000000	0.91022819	0.4647412
tax	-0.29204783	0.50645559	-0.53443158	0.910228189	1.00000000	0.4608530
ptratio	-0.35550149	0.26151501	-0.23247054	0.464741179	0.46085304	1.0000000
black	0.12806864	-0.27353398	0.29151167	-0.444412816	-0.44180801	-0.1773833
lstat	-0.61380827	0.60233853	-0.49699583	0.488676335	0.54399341	0.3740443
medv	0.69535995	-0.37695457	0.24992873	-0.381626231	-0.46853593	-0.5077867

	black	lstat	medv
crim	-0.38506394	0.4556215	-0.3883046
zn	0.17552032	-0.4129946	0.3604453
indus	-0.35697654	0.6037997	-0.4837252
chas	0.04878848	-0.0539293	0.1752602
nox	-0.38005064	0.5908789	-0.4273208
rm	0.12806864	-0.6138083	0.6953599
age	-0.27353398	0.6023385	-0.3769546
dis	0.29151167	-0.4969958	0.2499287
rad	-0.44441282	0.4886763	-0.3816262
tax	-0.44180801	0.5439934	-0.4685359
ptratio	-0.17738330	0.3740443	-0.5077867
black	1.00000000	-0.3660869	0.3334608
lstat	-0.36608690	1.0000000	-0.7376627
medv	0.33346082	-0.7376627	1.0000000

If we look at the median value of owner occupied homes in $1000 (medv) and try to understand its correlation between other variables then we can see that the average number of rooms (rm) has the highest positive correlation with medv and pupil/teacher ratio (ptratio) and % lower-status pop. (lstat) has high negative correlation with medv.

Regression Model

```
> RegModel.2 <- lm(medv~black+chas+crim+dis+lstat+nox+ptratio+rad+rm+zn,
+   data=Boston)

> summary(RegModel.2)

Call:
lm(formula = medv ~ black + chas + crim + dis + lstat + nox +
    ptratio + rad + rm + zn, data = Boston)

Residuals:
     Min       1Q   Median       3Q      Max
-16.2609  -2.9888  -0.5083   1.8041  26.2482  (The Residuals are quite high ranges from -16.26 to 26.24, we will have to bring this down)

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  34.712342   5.102742   6.803 2.97e-11 ***
black         0.009700   0.002701   3.591 0.000363 ***
chas          2.967868   0.860830   3.448 0.000614 ***
crim         -0.104843   0.033132  -3.164 0.001650 **
dis          -1.429370   0.186922  -7.647 1.08e-13 ***
lstat        -0.528147   0.047930 -11.019  < 2e-16 ***
nox         -20.314416   3.472292  -5.850 8.92e-09 ***
ptratio      -1.014914   0.129006  -7.867 2.30e-14 ***
rad           0.128761   0.040788   3.157 0.001692 **
rm            3.977104   0.407731   9.754  < 2e-16 ***
zn            0.036634   0.013412   2.731 0.006532 **
Signif. Codes:  0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1

Residual standard error: 4.79 on 495 degrees of freedom
Multiple R-squared:  0.7342,          Adjusted R-squared:  0.7288
F-statistic: 136.7 on 10 and 495 DF,  p-value: < 2.2e-16

> vif(RegModel.2)
   black     chas     crim      dis    lstat      nox  ptratio      rad
1.338982 1.052428 1.787963 3.410587 2.579040 3.564036 1.717222 2.776775
      rm       zn
1.806735 2.154054
```

The Variance Inflation Factors are fine because all the variables are below 4.

```
> outlierTest(RegModel.2)
    rstudent unadjusted p-value Bonferonni p
369 5.825701        1.0263e-08   5.1932e-06
372 5.394982        1.0650e-07   5.3886e-05
```

373 5.205160 2.8478e-07 1.4410e-04

There are 3 outliers we will now remove these outliers from the data and re-run the model

> Boston <- Boston[-c(369,372,373),]

> RegModel.2 <- lm(medv~black+chas+crim+dis+lstat+nox+ptratio+rad+rm+zn,
+ data=Boston)

> summary(RegModel.2)

Call:
lm(formula = medv ~ black + chas + crim + dis + lstat + nox +
ptratio + rad + rm + zn, data = Boston)

Residuals:
Min 1Q Median 3Q Max
-15.6641 -2.8410 -0.3751 1.8066 19.2247 (The range of the Residuals have come down considerably)

Coefficients:
Estimate Std. Error t value Pr(>|t|)
(Intercept) 28.640030 4.680205 6.119 1.92e-09 ***
black 0.009106 0.002454 3.711 0.00023 ***
chas 2.405876 0.792192 3.037 0.00252 **
crim -0.095376 0.030100 -3.169 0.00163 **
dis -1.229979 0.170831 -7.200 2.27e-12 ***
lstat -0.433398 0.044591 -9.719 < 2e-16 ***
nox -19.218002 3.154500 -6.092 2.25e-09 ***
ptratio -1.025144 0.117171 -8.749 < 2e-16 ***
rad 0.073739 0.037418 1.971 0.04932 *
rm 4.668916 0.377795 12.358 < 2e-16 ***
zn 0.030425 0.012193 2.495 0.01291 *

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 4.349 on 492 degrees of freedom
Multiple R-squared: 0.7699, Adjusted R-squared: 0.7653
F-statistic: 164.6 on 10 and 492 DF, p-value: < 2.2e-16

We will create another model eliminating the insignificant variables such as zn and rad

> RegModel.3 <- lm(medv~black+chas+crim+dis+lstat+nox+ptratio+rm, data=Boston)

> summary(RegModel.3)

Call:
lm(formula = medv ~ black + chas + crim + dis + lstat + nox +
ptratio + rm, data = Boston)

Residuals:
Min 1Q Median 3Q Max
-15.4796 -2.6529 -0.6067 1.6867 20.2662 (The range of the Residuals have slightly increased but still acceptable)

```
Coefficients:
             Estimate Std. Error t value Pr(>|t|)
(Intercept)  25.706692   4.450328   5.776 1.35e-08 ***
black         0.008151   0.002442   3.338 0.000907 ***
chas          2.382992   0.799459   2.981 0.003017 **
crim         -0.059040   0.027282  -2.164 0.030936 *
dis          -0.991221   0.149398  -6.635 8.56e-11 ***
lstat        -0.425606   0.044940  -9.470  < 2e-16 ***
nox         -16.937158   2.937265  -5.766 1.43e-08 ***
ptratio      -1.014710   0.101895  -9.958  < 2e-16 ***
rm            4.943351   0.372196  13.282  < 2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 4.39 on 494 degrees of freedom
Multiple R-squared:  0.7646,          Adjusted R-squared:  0.7608
F-statistic: 200.5 on 8 and 494 DF,  p-value: < 2.2e-16
```

> vif(RegModel.3)

```
   black     chas     crim      dis    lstat      nox  ptratio       rm
1.302030 1.051365 1.440841 2.569929 2.686973 3.024642 1.270217 1.778336
```

The Variance Inflation Factors are fine because all the variables are below 4.

Residual Analysis

Residual analysis is usually done graphically using:

- **Quantile plots: to assess normality**
- **Histograms and boxplots**
- **Scatterplots: to assess model assumptions, such as constant variance and linearity, and to identify potential outliers**

```
> mean(RegModel.3$residuals)
[1] 2.127881e-17

hist(RegModel.3$residuals, xlab="residuals", main="Histogram of residuals")

qqnorm(RegModel.3$residuals, main="Normal Probability Plot", pch=19)

qqline(RegModel.3$residuals)
```

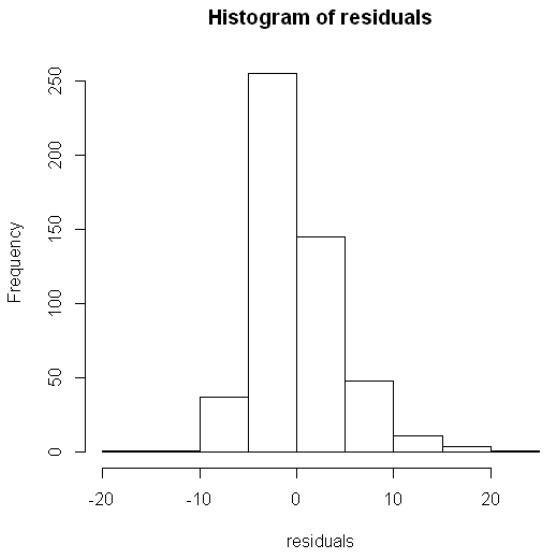

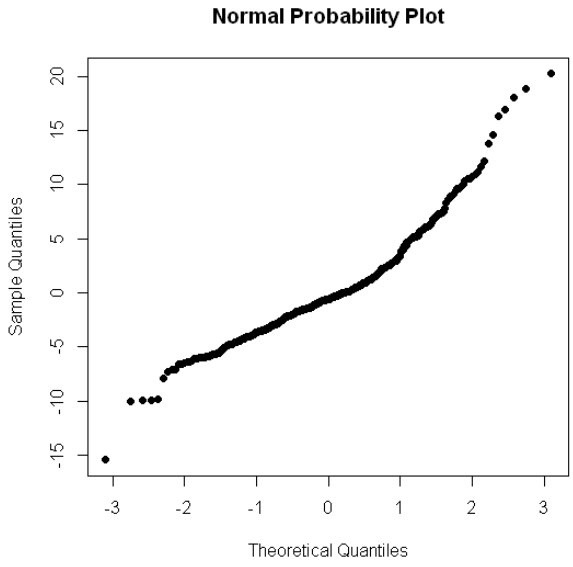

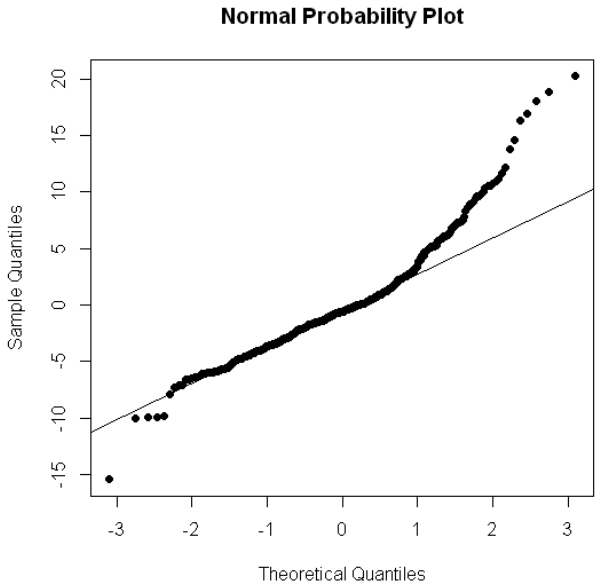

Checking for linear relationship amongst the important variables, error should have a constant variance and error terms should not be independent

Plotting residuals against key predictor variables, lstat, ptratio and rm

```
plot(lstat, RegModel.3$residuals, main="Residuals vs. Predictor", xlab="% in Lower Status", ylab="Residuals", pch=19)
abline(h=0)

plot(rm, RegModel.3$residuals, main="Residuals vs. Predictor", xlab="Average No of Rooms", ylab="Residuals", pch=19)
```

```
abline(h=0)

plot(ptratio, RegModel.3$residuals, main="Residuals vs. Predictor", xlab="Pupil/Teacher Ratio",
ylab="Residuals", pch=19)
abline(h=0)
```

Plotting residuals against fitted values

```
plot(RegModel.3$fitted.values, RegModel.3$residuals, main="Residuals vs. Fitted", xlab="Fitted values",
ylab="Residuals", pch=19)
abline(h=0)
```

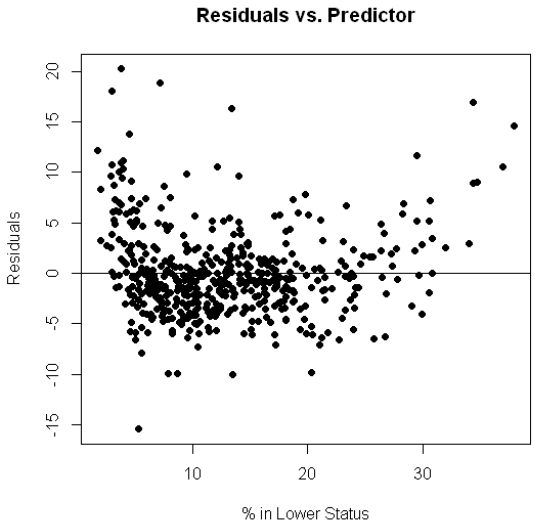

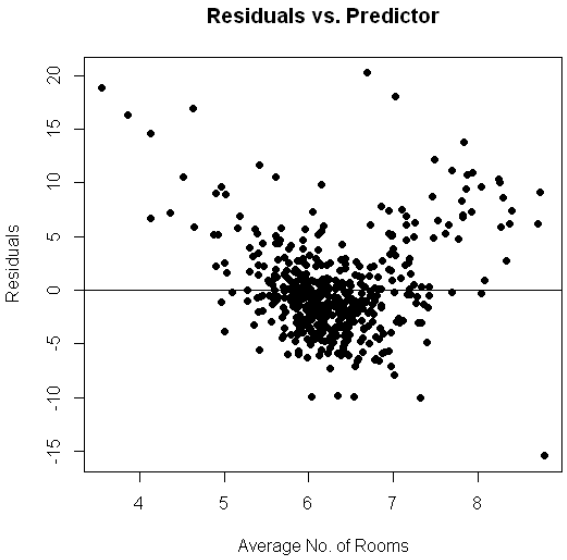

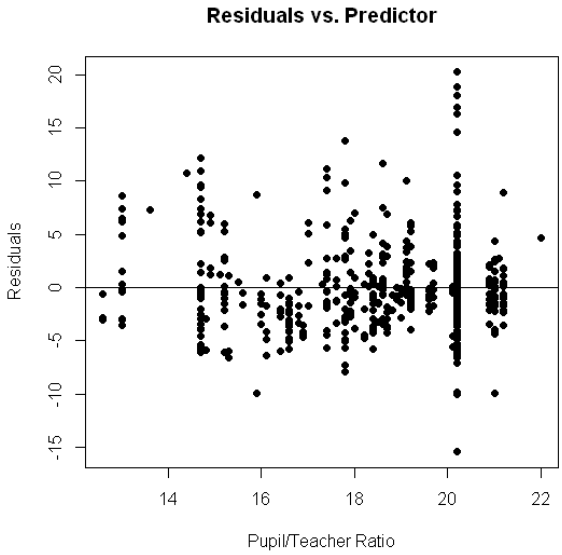

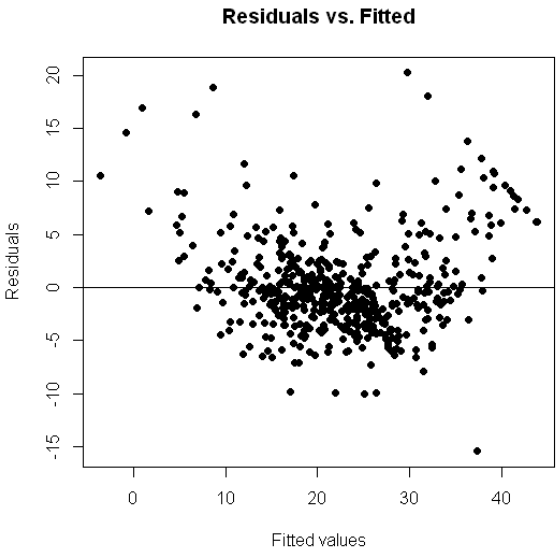

```
> anova(RegModel.2,RegModel.3)
```

Analysis of Variance Table

```
Model 1: medv ~ black + chas + crim + dis + lstat + nox + ptratio + rad +
    rm + zn
Model 2: medv ~ black + chas + crim + dis + lstat + nox + ptratio + rm
  Res.Df    RSS Df Sum of Sq      F    Pr(>F)
1    492 9303.9
2    494 9520.3 -2    -216.4 5.7216 0.003496 **
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```